

CS 492: Diffusion Models and its applications

Ayush Raina

June 8, 2025

Course Instructor: Prof. Minhyuk Sung

Semester: 7th

Department: Computer Science and Engineering

Contents

1. **Lecture 1:** Course Logistics
2. **Lecture 2:** Introduction to Generative Modelling

Lecture 1

Following is the comparison of DDPM, GAN and VAEs etc:

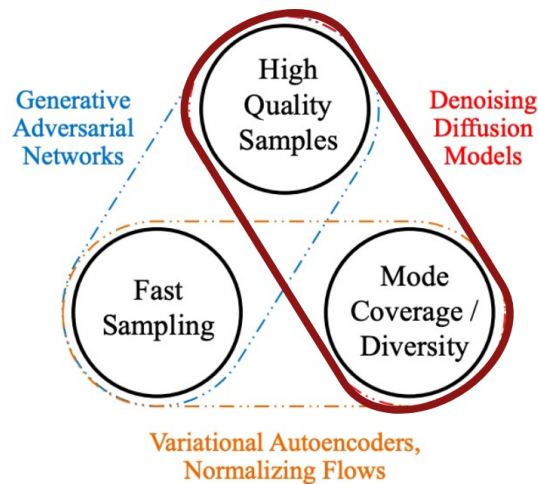


Figure 1: Comparison of DDPM, GAN, and VAEs

Diffusion Models generate high quality samples, cover diversity of data but sampling process is slow as compared to other generative models like GANs.

Following topics will be covered in this course:

1. Introduction to Generative Modelling
2. DDPM and Score Based Models
3. DDIM, Guided Diffusion and Latent Diffusion
4. Control Net, Conditional Generation and Personalization
5. Inverse Problem / Knowledge Distillation
6. Flow Based Models
7. Diffusion Transformers

There will be 7 assignments based on following topics:

1. DDPM
2. DDIM and CFG
3. Control Net and LORA
4. Distillation
5. Synchronization
6. DPM Solvers
7. Flow Based Models

Lecture 2

Sampling from a distribution

To sample from a discrete distribution, we need access to PMF. Suppose we have PMF p , such that $P(X = x_i) = p_i$ for $i = 1, 2, \dots, n$, $p_i \geq 0$ and $\sum_{i=1}^n p_i = 1$. Then we can use inverse transform sampling method to sample from this distribution.

Sampling Technique 1: Inverse Transform Sampling

First we calculate the CDF F of the PMF p as follows:

$$F(x_i) = P(X \leq x_i) = \sum_{j=1}^i p_j, F(x_n) = 1$$

then we can sample from the distribution as follows:

1. Sample $u \sim \mathcal{U}(0, 1)$.
2. Find i such that $F(x_{i-1}) \leq u \leq F(x_i)$.
3. Return x_i as the sample from the distribution.

In the case of continuous distributions, we can use the same method but with PDF instead of PMF. We can calculate the CDF as follows:

$$F(x) = P(X \leq x) = \int_{-\infty}^x p(t) dt$$

Then we can sample from the distribution as follows:

1. Sample $u \sim \mathcal{U}(0, 1)$.
2. Find x such that $F(x) = u$, but this involves the access to the inverse of CDF, which may not be available.
3. Return x as the sample from the distribution.

Example

Consider $p(x) = \frac{3}{8}x^2$ for $0 \leq x \leq 2$. Then we can calculate the CDF as follows:

$$F(x) = \int_0^x \frac{3}{8}t^2 dt = \frac{1}{8}x^3$$

In this case, we can calculate inverse easily and it turns out to be:

$$\begin{aligned} F^{-1}(u) &= (8u)^{1/3} \\ F^{-1}(u) &= 2(u)^{1/3} \end{aligned}$$

Hence sampling involves 2 steps here:

1. Sample $u \sim \mathcal{U}(0, 1)$.
2. Return $F^{-1}(u) = 2(u)^{1/3}$ as the sample from the distribution.

Sampling Technique 2: Rejection Sampling

If the inverse of the CDF cannot be computed, then:

1. Let $q(x)$ be a proposal (easy-to-sample) distribution and c be a constant such that $c \cdot q(x) \geq p(x)$ for all x .
2. Sample $x \sim q(x)$.
3. Sample $u \sim \mathcal{U}(0, 1)$.
4. Accept the sample x if $u \leq \frac{p(x)}{c \cdot q(x)}$; otherwise reject it and repeat from step 2.

Example: Rejection Sampling

Let's sample from the same distribution $p(x) = \frac{3}{8}x^2$ for $0 \leq x \leq 2$ using rejection sampling.

First, we need to find a proposal distribution $q(x)$ and a constant c such that $c \cdot q(x) \geq p(x)$ for all $x \in [0, 2]$.

Since $p(x) = \frac{3}{8}x^2$ is maximized at $x = 2$, we have $p(2) = \frac{3}{8} \cdot 4 = \frac{3}{2}$. We can choose $q(x) = \mathcal{U}(0, 2)$, which has PDF $q(x) = \frac{1}{2}$ for $x \in [0, 2]$. To ensure $c \cdot q(x) \geq p(x)$, choose $c = 3$, since $c \cdot q(x) = 3 \cdot \frac{1}{2} = \frac{3}{2}$, which matches the maximum of $p(x)$.

The rejection sampling steps are:

1. Sample $x \sim \mathcal{U}(0, 2)$.
2. Sample $u \sim \mathcal{U}(0, 1)$.
3. Accept x if $u \leq \frac{p(x)}{c \cdot q(x)} = \frac{\frac{3}{8}x^2}{3 \cdot \frac{1}{2}} = \frac{x^2}{4}$; otherwise, reject and repeat.

Generative Adversarial Networks (GANs)

Loss Function:

$$\min_G \max_D \mathcal{L}(D, G) = \mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$$

Where D is the discriminator, G is the generator, p_{data} is the data distribution, and p_z is the noise distribution. GAN training involves challenges like **mode collapse** and **non-convergence**.

Non Convergence

Discriminator becomes too good and easily classifies the real and fake samples and as a result leading to zero gradient flow for generator and hence no learning.

Mode Collapse

Generator produces samples from a single mode of the data distribution because for this mode the discriminator is not able to distinguish between real and fake samples. But this does not capture the diversity of the data distribution.

Variational Autoencoders (VAEs)

This is a generative model which does not involve solving the min-max optimization problem. KL Divergence between 2 distributions is defined as:

$$D_{KL}(p||q) = \int p(x) \log \frac{p(x)}{q(x)} dx = \mathbb{E}_{x \sim p(x)} \left[\log \frac{p(x)}{q(x)} \right]$$

Homework: If $p(x) = \mathcal{N}(x; \mu, \sigma^2 \mathbb{I})$ and $q(x) = \mathcal{N}(x; 0, \mathbb{I})$, then find the KL divergence $D_{KL}(p||q)$.

Convex Function

A function f is convex if for all x, y and $\lambda \in [0, 1]$:

$$f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y)$$

In general f is convex if $f(\sum_{i=1}^n \lambda_i x_i) \leq \sum_{i=1}^n \lambda_i f(x_i)$ for all x_i and $\lambda_i \geq 0$ such that $\sum_{i=1}^n \lambda_i = 1$.

Jensen's Inequality

If f is a convex function and X is a random variable, then:

$$\mathbb{E}_{p(x)}[f(X)] \geq f(\mathbb{E}_{p(x)}[X])$$

whereas for concave functions, the inequality is reversed:

$$\mathbb{E}_{p(x)}[f(X)] \leq f(\mathbb{E}_{p(x)}[X])$$

Latent Variable Models introduce hidden (unobserved) variables \mathbf{z} to explain observed data \mathbf{x} . The generative process is:

- Sample latent variable: $\mathbf{z} \sim p(\mathbf{z})$
- Generate data: $\mathbf{x} \sim p_\theta(\mathbf{x} | \mathbf{z})$

The marginal likelihood of data is:

$$p_\theta(\mathbf{x}) = \int p_\theta(\mathbf{x} | \mathbf{z}) p(\mathbf{z}) d\mathbf{z}$$

Difference: Previous models directly modeled $p_\theta(\mathbf{x})$; latent variable models introduce \mathbf{z} to capture hidden structure or factors. Learning involves maximizing the marginal likelihood:

$$\theta^* = \arg \max_{\theta} \sum_{i=1}^N \log \int p_\theta(\mathbf{x}_i | \mathbf{z}) p(\mathbf{z}) d\mathbf{z}$$

Problems

We can easily see that this is better modelling choice. We can try to find:

$$\begin{aligned}\theta^* &= \arg \max_{\theta} \sum_{i=1}^N \log p_{\theta}(\mathbf{x}_i) \\ &= \arg \max_{\theta} \sum_{i=1}^N \log \sum_{z_i \sim \mathcal{Z}} p_{\theta}(\mathbf{x}_i, z_i)\end{aligned}$$

But for a particular x_i : z_i 's are not observed. Hence we need to marginalize over all possible z_i 's which is not feasible. Hence we need some other approach to do MLE estimation on latent variable models.

Evidence Lower Bound (ELBO)

We saw that computing log likelihood of partially observed data $p_{\theta}(\mathbf{x}, z)$ is hard to compute. We will instead construct a lower bound on LL.

1. Jensen's Inequality: $f(\mathbb{E}[X]) \leq \mathbb{E}[f(X)]$ for convex function f , where as for concave function f we have $f(\mathbb{E}[X]) \geq \mathbb{E}[f(X)]$.

2. $\log(x)$ is a concave function. This means:

$$\begin{aligned}\log(p_{\theta}(\mathbf{x})) &= \log \left(\int p_{\theta}(\mathbf{x}, z) dz \right) = \log \left(\int q(z) \frac{p_{\theta}(\mathbf{x}, z)}{q(z)} dz \right) \\ &= \log \left(\mathbb{E}_{z \sim q(z)} \left[\frac{p_{\theta}(\mathbf{x}, z)}{q(z)} \right] \right) \geq \mathbb{E}_{z \sim q(z)} \left[\log \left(\frac{p_{\theta}(\mathbf{x}, z)}{q(z)} \right) \right] := ELBO\end{aligned}$$

$$\mathcal{L}_{ELBO}(\theta) := \mathbb{E}_{z \sim q(z)} \left[\log \left(\frac{p_{\theta}(\mathbf{x}, z)}{q(z)} \right) \right]$$

We have shown that ELBO is a lower bound on log likelihood and little more math shows that when $q(z) = p_{\theta}(z|\mathbf{x})$ then equality holds.

Variational Autoencoders (VAEs)

We approximate the posterior $p_{\theta}(z | \mathbf{x})$ with $q_{\phi}(z | \mathbf{x})$ (encoder neural network). Hence our loss function becomes:

$$\begin{aligned}\mathcal{L}_{VAE}(\theta, \phi) &= \mathbb{E}_{z \sim q_{\phi}(z|\mathbf{x})} \left[\log \left(\frac{p_{\theta}(\mathbf{x}, z)}{q_{\phi}(z | \mathbf{x})} \right) \right] \\ &= \mathbb{E}_{z \sim q_{\phi}(z|\mathbf{x})} [\log(p(z)) + \log(p_{\theta}(\mathbf{x} | z)) - \log(q_{\phi}(z | \mathbf{x}))] \\ &= \mathbb{E}_{z \sim q_{\phi}(z|\mathbf{x})} \left[\log(p_{\theta}(\mathbf{x} | z)) - \log \left(\frac{q_{\phi}(z | \mathbf{x})}{p(z)} \right) \right] \\ \mathcal{L}_{VAE}(\theta, \phi) &= \mathbb{E}_{z \sim q_{\phi}(z|\mathbf{x})} [\log(p_{\theta}(\mathbf{x} | z))] - D_{KL}(q_{\phi}(z | \mathbf{x}) || p(z))\end{aligned}$$

Here is the final loss function we need to optimize:

$$\mathcal{L}_{VAE}(\theta, \phi) = \mathbb{E}_{z \sim q_{\phi}(z|\mathbf{x})} [\log(p_{\theta}(\mathbf{x} | z))] - D_{KL}(q_{\phi}(z | \mathbf{x}) || p(z))$$

1. In VAE literature, $q_{\phi}(z | \mathbf{x})$ is called the **encoder** and $p_{\theta}(\mathbf{x} | z)$ is called the **decoder**.

2. The first term is the **reconstruction loss** (how well the model can reconstruct the data given the latent variable).

3. The second term is the **KL divergence** between the approximate posterior and the prior. It regularizes the latent space to follow a specific distribution (usually Gaussian).

We choose $p(z)$ standard normal distribution $\mathcal{N}(0, I)$ and $q_{\phi}(z | \mathbf{x}) = \mathcal{N}(\mathbf{x}; \mu_{\phi}(\mathbf{x}), \sigma_{\phi}(\mathbf{x}))$ and $p_{\theta}(\mathbf{x} | z) = \mathcal{N}(\mathbf{x}; \mu_{\theta}(z), \sigma_{\theta}(z))$.